

## **Spotify Dataset Analysis**



“ Who needs therapy when you have Spotify 😊”

Group 15

Nihar Khillar

Parveen Mittal

Sahejbir Singh Kumar

Simran Sharma

## TABLE OF CONTENTS

<b><u>LINK TO POWERPOINT PRESENTATION RECORDING .....</u></b>	<b><u>3</u></b>
<b><u>EXECUTIVE SUMMARY .....</u></b>	<b><u>3</u></b>
<b><u>PROJECT MOTIVATION/ BACKGROUND .....</u></b>	<b><u>4</u></b>
<b><u>PROJECT FLOW .....</u></b>	<b><u>4</u></b>
<b><u>DATA DESCRIPTION.....</u></b>	<b><u>4</u></b>
<b><u>DATA CLEANING AND PRE PROCESSING .....</u></b>	<b><u>6</u></b>
<b><u>DATA CLEANING .....</u></b>	<b><u>7</u></b>
<b><u>DATA PREPROCESSING .....</u></b>	<b><u>9</u></b>
<b><u>EXPLORATORY DATA ANALYSIS .....</u></b>	<b><u>11</u></b>
<b><u>DATA MODELING .....</u></b>	<b><u>20</u></b>
<b><u>DECISION TREE MODELS .....</u></b>	<b><u>20</u></b>
<b><u>LOGISTIC REGRESSION MODELS .....</u></b>	<b><u>26</u></b>
<b><u>FINAL MODEL SELECTION .....</u></b>	<b><u>30</u></b>
<b><u>CLUSTERING WITH K MEANS.....</u></b>	<b><u>30</u></b>
<b><u>KEY TAKEAWAYS.....</u></b>	<b><u>33</u></b>

## LINK TO POWERPOINT PRESENTATION RECORDING

[https://cometmail-my.sharepoint.com/:v:/g/personal/pxk220018\\_utdallas\\_edu/EcxPXbxPriJliBuVv2IGMjkBi1x6\\_Ydwnw1pJOqSZH01oQ?e=2cPAXg](https://cometmail-my.sharepoint.com/:v:/g/personal/pxk220018_utdallas_edu/EcxPXbxPriJliBuVv2IGMjkBi1x6_Ydwnw1pJOqSZH01oQ?e=2cPAXg)

## EXECUTIVE SUMMARY

Music is something that is close to each one of us. Be it tough times, be it easy times, be it happy or be it sad times, there's music for each of those occasions. Still, we know so little about music. We all have our favorite songs, podcast or audio book but we fail to describe perfectly why do we like them. Through this project, we try to do so through the analysis of the Spotify data set.

We had set 3 clear objectives for this project:

- Whether music features impact the popularity of a track: The business relevance is music artists can focus on the key features in future projects to increase popularity
- Build a model to predict the popularity of a track: Business relevance: Music producers and Spotify can select music compositions with high probability of popularity
- Make clusters to club similar tracks together based on music features: Business relevance: Can be used to suggest songs to users based on their past listening record

We followed the below process for carrying out the project:

1. Data collection: secondhand data set was collected from Kaggle
2. Data cleaning and preprocessing: In this step, we focused on 3 objectives:
  - a. Finding the appropriate subset of data from the large dataset since the original dataset had over 580,000 rows of data. We did so by selecting the data of songs released between 16<sup>th</sup> April 2020 to 16<sup>th</sup> April 2021
  - b. Getting rid of the duplicate data and rows with missing values (NA or blank spaces)
  - c. Convert the numerical popularity score (0-100) to a categorical variable popularity category which will be the dependent variable going forward
3. Exploratory data analysis: In this step, we focused on using graphical tools and statistical summaries available in R to get a better sense of the data as well as identifying key patterns and features in the data set
4. Data Modeling: We focused on three tasks in this step:
  - a. Through decision tree and logistic regression models, understand the key music features which explain most of the track's popularity
  - b. Build multiple models which can predict the track's popularity based on the relevant input features and test these models with the validation dataset
  - c. Through clustering, create clusters of similar tracks (K means) which can act as a recommendation to listeners based on their music preferences
5. Sharing our insights and key takeaways

## PROJECT MOTIVATION/ BACKGROUND

- There has always been a close connection between music and society. The music both produces and reflects societal conditions, particularly those that either support or resist social change. The way that the majority of people receive music has changed dramatically since the introduction of recording techniques in the second part of the 20th century. The majority of people have instant access to all genres of music, day and night.
- Music can help us sleep, get pumped up for school and work, calm us down after a bad day, keep our spirits up when we're feeling low, and encourage socialization between people.
- With easy access to smart applications and the creation of music according to genres liked by the people, the music industry is dominating the various industries by providing society with a source of entertainment and stress-free days.
- When talking about online music streaming, one cannot ignore the contributions of Spotify to the music industry. It has attracted millions of subscribers to try and pay for the music they listen to, reducing piracy substantially.
- A huge library of tracks, a great experience across all connected devices, playlist creation tools, and befitting personalised recommendations make Spotify stand out and give it a monthly active user base of 381 million.<sup>1</sup>
- With availability in more than 170 countries<sup>2</sup> at the moment, the Spotify business model has the upper hand over many others because of the significant number of features provided and the delivery of music with no delay.
- The humongous use of the Spotify application led us to choose this domain for our analysis and derive some interesting insights by building the prediction model.

## PROJECT FLOW

The R codes have been presented along with key outputs and our observations for each process of the data analysis process

## DATA DESCRIPTION

The second hand dataset was sourced from Kaggle. Link is given below.

(<https://www.kaggle.com/datasets/lehaknarnauli/spotify-datasets?select=tracks.csv>)

The following libraries were used in the analysis in R: moments, ggplot2, GGally, rpart, rpart.plot, caret, dplyr, funModeling, randomForest, adabag

```
spotify.df <- read.csv("Spotify.csv")
```

```
View(spotify.df)
```

---

<sup>1</sup> <https://www.feedough.com/how-does-spotify-make-money/>

<sup>2</sup> <https://www.feedough.com/how-does-spotify-make-money/>

## #Understanding the dataset

```
summary(spotify.df)
```

```
str(spotify.df)
```

```
sapply(colnames(spotify.df), function(x) class(spotify.df[[x]]))
```

```
lapply(spotify.df, unique)
```

```
View(spotify.df)
```

```
dim(spotify.df)
```

```
      id      name      popularity      duration_ms      explicit
Length:586672 Length:586672 Min.   : 0.00 Min.   : 3344 Min.   :0.00000
Class :character Class :character 1st Qu.: 13.00 1st Qu.: 175093 1st Qu.:0.00000
Mode  :character Mode  :character Median : 27.00 Median : 214893 Median :0.00000
                                Mean  : 27.57 Mean  : 230051 Mean  :0.04409
                                3rd Qu.: 41.00 3rd Qu.: 263867 3rd Qu.:0.00000
                                Max.   :100.00 Max.   :5621218 Max.   :1.00000

      artists      id_artists      release_date      danceability      energy
Length:586672 Length:586672 Length:586672 Min.   :0.0000 Min.   :0.000
Class :character Class :character Class :character 1st Qu.:0.4530 1st Qu.:0.343
Mode  :character Mode  :character Mode  :character Median :0.5770 Median :0.549
                                Mean  :0.5636 Mean  :0.542
                                3rd Qu.:0.6860 3rd Qu.:0.748
                                Max.   :0.9910 Max.   :1.000

      key      loudness      mode      speechiness      acousticness
Min.   : 0.000 Min.   : -60.000 Min.   :0.00000 Min.   :0.0000 Min.   :0.0000
1st Qu.: 2.000 1st Qu.: -12.891 1st Qu.:0.00000 1st Qu.:0.0340 1st Qu.:0.0969
Median : 5.000 Median : -9.243 Median :1.00000 Median :0.0443 Median :0.4220
Mean   : 5.222 Mean   : -10.206 Mean   :0.6588 Mean   :0.1049 Mean   :0.4499
3rd Qu.: 8.000 3rd Qu.: -6.482 3rd Qu.:1.00000 3rd Qu.:0.0763 3rd Qu.:0.7850
Max.   :11.000 Max.   : 5.376 Max.   :1.00000 Max.   :0.9710 Max.   :0.9960

      instrumentalness      liveness      valence      tempo      time_signature
Min.   :0.000000000 Min.   :0.0000 Min.   :0.00000 Min.   : 0.0 Min.   :0.000
1st Qu.:0.000000000 1st Qu.:0.0983 1st Qu.:0.3460 1st Qu.: 95.6 1st Qu.:4.000
Median :0.0000245 Median :0.1390 Median :0.5640 Median :117.4 Median :4.000
Mean   :0.1134508 Mean   :0.2139 Mean   :0.5523 Mean   :118.5 Mean   :3.873
3rd Qu.:0.0095500 3rd Qu.:0.2780 3rd Qu.:0.7690 3rd Qu.:136.3 3rd Qu.:4.000
Max.   :1.0000000 Max.   :1.0000 Max.   :1.0000 Max.   :246.4 Max.   :5.000

      id      name      popularity      duration_ms      explicit
"character" "character" "integer" "integer" "integer"
      artists      id_artists      release_date      danceability      energy
"character" "character" "character" "numeric" "numeric"
      key      loudness      mode      speechiness      acousticness
"integer" "numeric" "integer" "numeric" "numeric"
instrumentalness      liveness      valence      tempo      time_signature
"numeric" "numeric" "numeric" "numeric" "integer"
```

There were 586672 rows and 20 columns in the data set. There were 3 major categorical variables which were of interest in the analysis: mode, key and explicit while other numerical features of importance were duration\_ms, danceability, energy, loudness, mode, speechiness,

acousticness, instrumentalness, liveness, valence, tempo and time\_signature. The brief description of these features are as under

- Duration\_ms: Duration of track in milliseconds
- Explicit: Whether the track has explicit lyrics (1: Y, 0: N)
- Danceability: Track suitability for dancing
- Energy: Perceptual measure of intensity and activity
- Key: Overall key of the track
- Loudness: Overall loudness in decibels (dB)
- Mode: Depicts the modality (major 1 or minor 0) of a track
- Speechiness: Measure of presence of spoken words in a track
- Acousticness: Whether the track is acoustic.
- Instrumentalness: Measures whether a track contains no vocals
- Liveness: Detects the presence of live audience in the recording
- Valence: Describes the musical positiveness conveyed by a track
- Tempo: overall estimated tempo of a track in beats per minute
- Time\_signature: stimated overall time signature of a track
- Popularity: Measure the popularity from 0 to 100 based on number of times the tracks have been played

## DATA CLEANING AND PRE PROCESSING

The latest release date of the track in the dataset was 16<sup>th</sup> April 2021. We decided to include a year data i.e. between 16<sup>th</sup> April 2020 and 16<sup>th</sup> April 2021 for this analysis to avoid running into system crashes in the later steps. The dataset was reduced to 17328 rows now. Initial dataset had 586672 rows as mentioned earlier.

```
spotify2.df<-spotify.df[spotify.df$release_date >= "2020-04-16" & spotify.df$release_date <= "2021-04-16", ]
```

#Understanding the trimmed dataset

```
summary(spotify2.df)
```

```
str(spotify2.df)
```

```
lapply(spotify2.df, unique)
```

```
head(spotify2.df)
```

```
glimpse(spotify2.df)
```

```
colnames(spotify2.df)
```

```
dim(spotify2.df)
```

## DATA CLEANING

#Checking for NA values in the data set

```
colSums(is.na(spotify2.df))
```

#Checking for blank spaces in the numeric/ integer columns

```
nrow(spotify2.df[spotify2.df$popularity=="",])
```

```
nrow(spotify2.df[spotify2.df$duration_ms=="",])
```

```
nrow(spotify2.df[spotify2.df$explicit=="",])
```

```
nrow(spotify2.df[spotify2.df$danceability=="",])
```

```
nrow(spotify2.df[spotify2.df$energy=="",])
```

```
nrow(spotify2.df[spotify2.df$key=="",])
```

```
nrow(spotify2.df[spotify2.df$mode=="",])
```

```
nrow(spotify2.df[spotify2.df$speechiness=="",])
```

```
nrow(spotify2.df[spotify2.df$acousticness=="",])
```

```
nrow(spotify2.df[spotify2.df$instrumentalness=="",])
```

```
nrow(spotify2.df[spotify2.df$liveness=="",])
```

```
nrow(spotify2.df[spotify2.df$valence=="",])
```

```
nrow(spotify2.df[spotify2.df$tempo=="",])
```

```
nrow(spotify2.df[spotify2.df$time_signature=="",])
```

#Checking for duplicated values

```
sum(duplicated(spotify2.df))
```

#Counting unique songs in the data set

```
length(unique(spotify2.df$id))
```

```
> colSums(is.na(spotify2.df)) # No NA values
```

id	name	popularity	duration_ms	explicit
0	0	0	0	0
artists	id_artists	release_date	danceability	energy
0	0	0	0	0
key	loudness	mode	speechiness	acousticness
0	0	0	0	0
instrumentalness	liveness	valence	tempo	time_signature
0	0	0	0	0
popularity_cat				
0				

```

> nrow(spotify2.df[spotify2.df$popularity=="",])
[1] 0
> nrow(spotify2.df[spotify2.df$duration_ms=="",])
[1] 0
> nrow(spotify2.df[spotify2.df$explicit=="",])
[1] 0
> nrow(spotify2.df[spotify2.df$danceability=="",])
[1] 0
> nrow(spotify2.df[spotify2.df$energy=="",])
[1] 0
> nrow(spotify2.df[spotify2.df$key=="",])
[1] 0
> nrow(spotify2.df[spotify2.df$mode=="",])
[1] 0
> nrow(spotify2.df[spotify2.df$speechiness=="",])
[1] 0
> nrow(spotify2.df[spotify2.df$acousticness=="",])
[1] 0
> nrow(spotify2.df[spotify2.df$instrumentalness=="",])
[1] 0
> nrow(spotify2.df[spotify2.df$liveness=="",])
[1] 0
> nrow(spotify2.df[spotify2.df$valence=="",])
[1] 0
> nrow(spotify2.df[spotify2.df$tempo=="",])
[1] 0
> nrow(spotify2.df[spotify2.df$time_signature=="",])
[1] 0

> sum(duplicated(spotify2.df))
[1] 0
> length(unique(spotify2.df$id)) #counting unique songs
[1] 17328

```


There were no NAs or blank spaces and no duplicate data in the dataset.



## DATA PREPROCESSING

```
hist(spotify2.df$popularity, breaks=20, xlab = "popularity")
```

**Histogram of spotify2.df\$popularity**



In the above histogram, we observed that there were more than 3000 songs in the data set with popularity score in the range of 0 to 5. We decided to remove the songs with 0 popularity score. We wanted to analyze the songs which had been listened to at least once. We replaced the 0 values with NA first and then removed the NA value rows. About 1940 rows were removed leaving us with the final dataset of 15388 rows.

```
spotify2.df[\"popularity\"][spotify2.df[\"popularity\"] == 0] <- NA
```

```
summary(spotify2.df$popularity)
```

```
spotify2.df <- spotify2.df[complete.cases(spotify2.df),]
```

```
dim(spotify2.df)
```


```
> summary(spotify2.df$popularity)
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   NA's
  1.00  37.00   52.00  46.54  63.00  100.00  1940
> spotify2.df <- spotify2.df[complete.cases(spotify2.df),]
> dim(spotify2.df)
[1] 15388    20
```

Using mutate, we would add popularity\_cat in the dataset based on the popularity score. We decided that a song would be a hit if it had a popularity score of over 50 and not hit if score is less than equal to 50.

```
spotify2.df <- spotify2.df %>% mutate(popularity_cat = case_when(popularity <= 50 ~ 0, popularity > 50 ~ 1))
```

```
ggplot(spotify2.df, aes(factor(popularity_cat)))+geom_bar()+xlab("popularity_cat")
```

```
table(spotify2.df$popularity_cat)
```



```
> table(spotify2.df$popularity_cat)
```

```
  0    1
7282 8106
```

There were 7282 not hit songs and 8106 hit songs in the dataset.

## EXPLORATORY DATA ANALYSIS

```
numerical_cols = c("duration_ms", "danceability", "energy", "loudness", "speechiness", "acousticness",
                    "instrumentalness", "liveness", "valence", "tempo", "popularity")
```

```
spotify2_num.df <- spotify2.df %>% select(numerical_cols)
```

```
data.frame(mean=sapply(spotify2_num.df, mean, na.rm=TRUE),
```

```
          median=sapply(spotify2_num.df, median, na.rm=TRUE),
```

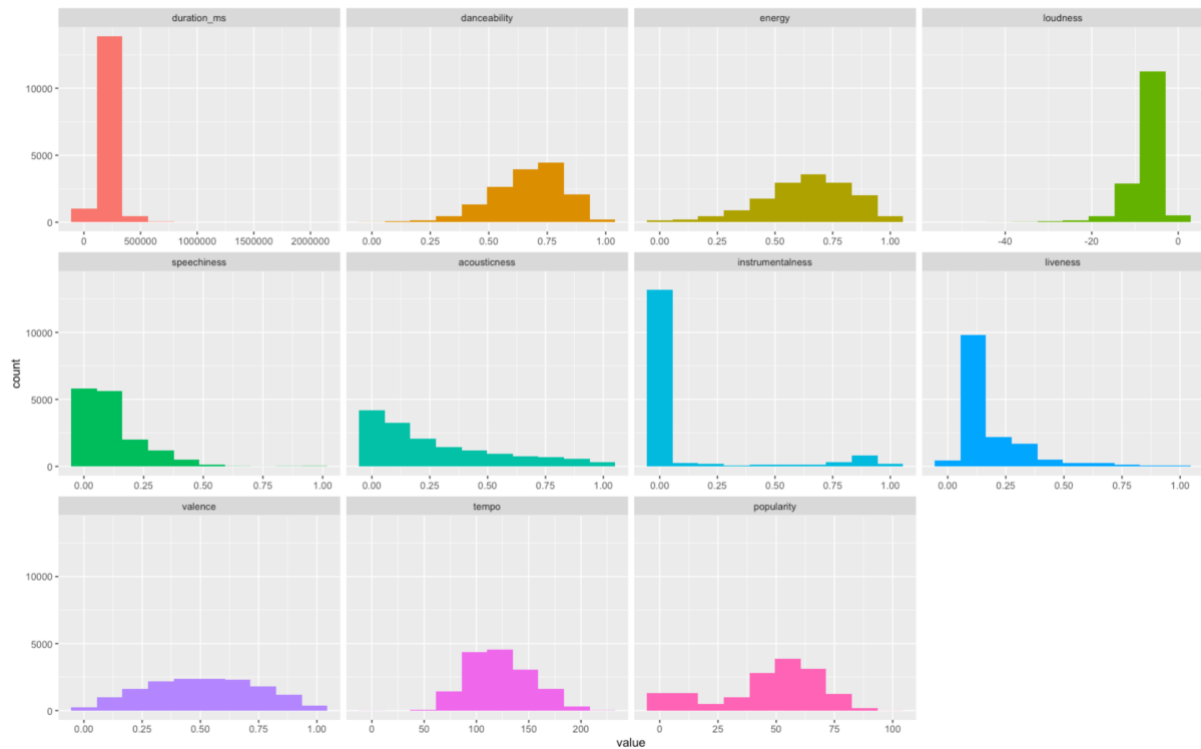
```
          min=sapply(spotify2_num.df, min, na.rm=TRUE),
```

```
          max=sapply(spotify2_num.df, max, na.rm=TRUE),
```

```
          standardev=sapply(spotify2_num.df, sd, na.rm=TRUE))
```

	mean	median	min	max	standardev
duration_ms	197493.77410970	192650.00000000	23534.0	2059336.000	69572.7428111
danceability	0.66939446	0.68800000	0.0	0.986	0.1517994
energy	0.63415982	0.65100000	0.0	1.000	0.1960134
loudness	-7.49118339	-6.71600000	-51.8	1.509	3.8526665
speechiness	0.12267376	0.07020000	0.0	0.966	0.1195391
acousticness	0.27290187	0.17700000	0.0	0.996	0.2708028
instrumentalness	0.09731392	0.00000133	0.0	1.000	0.2605267
liveness	0.18146670	0.12200000	0.0	0.992	0.1476501
valence	0.50822166	0.51000000	0.0	0.990	0.2341557
tempo	122.04801118	122.00050000	0.0	220.470	29.1366760
popularity	46.53782168	52.00000000	1.0	100.000	22.2820925

Duration\_ms had a very high standard deviation and high range (Max values were for yearbook and new years mix which tend to club multiple tracks together which increased the length of the tracks)



#### #Skewness check

```
skewness(spotify2_num.df$duration_ms,na.rm=TRUE)
```

```
skewness(spotify2_num.df$danceability,na.rm=TRUE)
```

```
skewness(spotify2_num.df$energy,na.rm=TRUE)
```

```
skewness(spotify2_num.df$loudness,na.rm=TRUE)
```

```
skewness(spotify2_num.df$speechiness,na.rm=TRUE)
```

```
skewness(spotify2_num.df$acousticness,na.rm=TRUE)
```

```
skewness(spotify2_num.df$instrumentalness,na.rm=TRUE)
```

```
skewness(spotify2_num.df$liveness,na.rm=TRUE)
```

```
skewness(spotify2_num.df$valence,na.rm=TRUE)
```

```
skewness(spotify2_num.df$tempo,na.rm=TRUE)
```

```
skewness(spotify2_num.df$popularity,na.rm=TRUE)
```

Based on histograms observation and from skewness metric, comments on the skewness of the data frame variables are given below:

- Almost symmetric: Valence, tempo
- Moderately skewed: Acousticness, popularity, energy, danceability
- Highly skewed: duration\_ms, loudness, speechiness, liveness, instrumentalness

We checked the spread of data for popularity\_cat across multiple features.

```
#boxplot
```

```
boxplot(duration_ms~popularity_cat, data = spotify2.df,
```

```
    main = "Variation",
```

```
    xlab = "duration_ms",
```

```
    ylab = "Popularity Category (0=Not Hit, 1=Hit)",
```

```
    col = "green",
```

```
    border = "blue",
```

```
    horizontal = TRUE,
```

```
    notch = TRUE
```

```
)
```

```
boxplot(danceability~popularity_cat, data = spotify2.df,
```

```
    main = "Variation",
```

```
    xlab = "danceability",
```

```
    ylab = "Popularity Category (0=Not Hit, 1=Hit)",
```

```
    col = "green",
```

```
    border = "blue",
```

```
    horizontal = TRUE,
```

```
    notch = TRUE
```

```
)
```

```
boxplot(energy~popularity_cat, data = spotify2.df,
```

```
    main = "Variation",
```

```
    xlab = "energy",
```

```
    ylab = "Popularity Category (0=Not Hit, 1=Hit)",
```

```
    col = "green",
```

```
    border = "blue",
```

```
    horizontal = TRUE,
```

```

        notch = TRUE
    )
boxplot(loudness~popularity_cat, data = spotify2.df,
        main = "Variation",
        xlab = "loudness",
        ylab = "Popularity Category (0=Not Hit, 1=Hit)",
        col = "green",
        border = "blue",
        horizontal = TRUE,
        notch = TRUE
    )
boxplot(speechiness~popularity_cat, data = spotify2.df,
        main = "Variation",
        xlab = "speechiness",
        ylab = "Popularity Category (0=Not Hit, 1=Hit)",
        col = "green",
        border = "blue",
        horizontal = TRUE,
        notch = TRUE
    )
boxplot(acousticness~popularity_cat, data = spotify2.df,
        main = "Variation",
        xlab = "acousticness",
        ylab = "Popularity Category (0=Not Hit, 1=Hit)",
        col = "green",
        border = "blue",
        horizontal = TRUE,

```

```


        notch = TRUE
    )
boxplot(instrumentalness~popularity_cat, data = spotify2.df,
        main = "Variation",
        xlab = "instrumentalness",
        ylab = "Popularity Category (0=Not Hit, 1=Hit)",
        col = "green",
        border = "blue",
        horizontal = TRUE,
        notch = TRUE
    )
boxplot(liveness~popularity_cat, data = spotify2.df,
        main = "Variation",
        xlab = "liveness",
        ylab = "Popularity Category (0=Not Hit, 1=Hit)",
        col = "green",
        border = "blue",
        horizontal = TRUE,
        notch = TRUE
    )
boxplot(valence~popularity_cat, data = spotify2.df,
        main = "Variation",
        xlab = "valence",
        ylab = "Popularity Category (0=Not Hit, 1=Hit)",
        col = "green",
        border = "blue",
        horizontal = TRUE,

```







```




    notch = TRUE
)
boxplot(tempo~popularity_cat, data = spotify2.df,
    main = "Variation",
    xlab = "tempo",
    ylab = "Popularity Category (0=Not Hit, 1=Hit)",
    col = "green",
    border = "blue",
    horizontal = TRUE,
    notch = TRUE
)
boxplot(popularity~popularity_cat, data = spotify2.df,
    main = "Variation",
    xlab = "popularity",
    ylab = "Popularity Category (0=Not Hit, 1=Hit)",
    col = "green",
    border = "blue",
    horizontal = TRUE,
    notch = TRUE
)

```



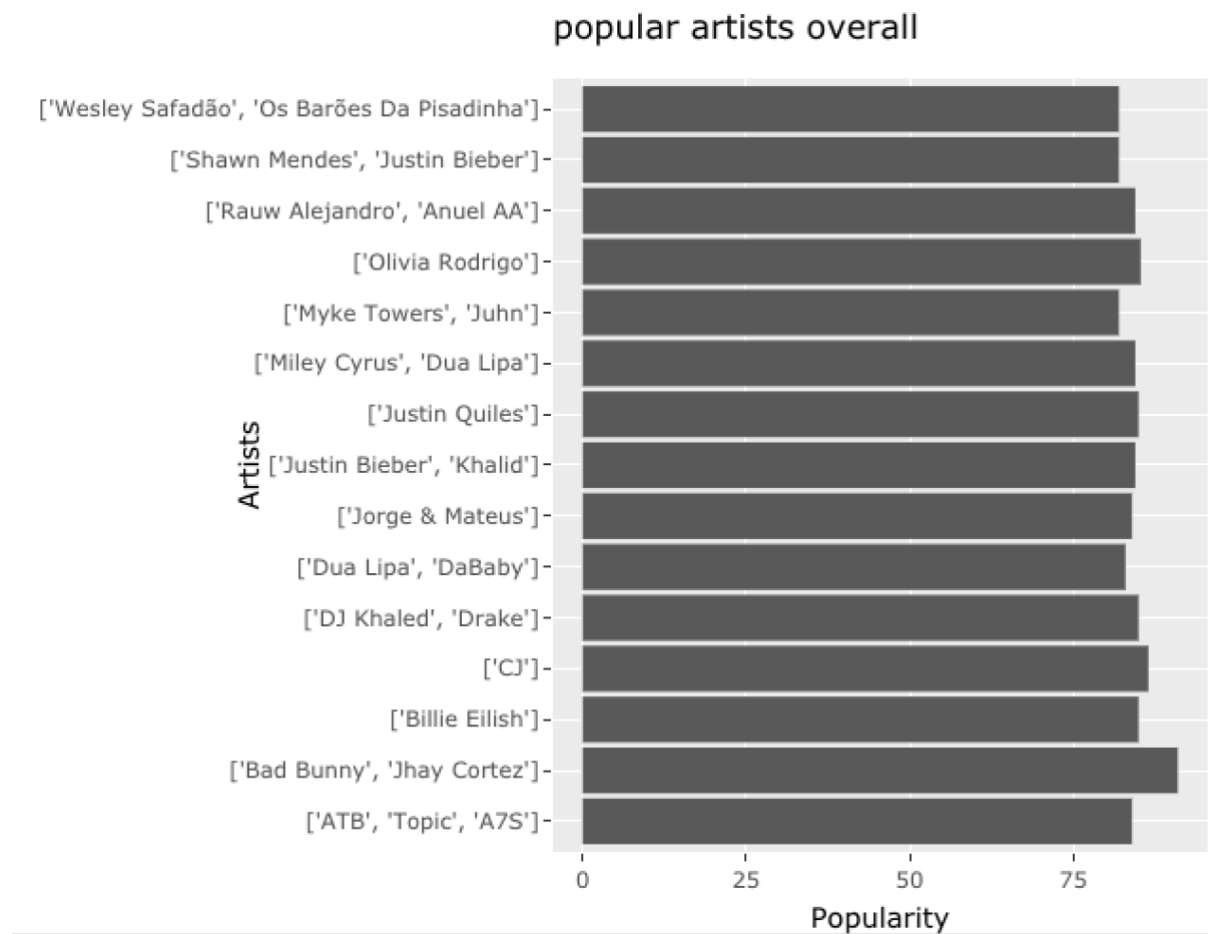






#finding popular artists

```
popular_artists <- spotify2.df %>% group_by(Artist = artists) %>%
  summarise(No_of_tracks = n(), Popularity = mean(popularity)) %>%
  filter(No_of_tracks > 1) %>%
  arrange(desc(Popularity)) %>%
  top_n(15, wt = Popularity) %>%
  ggplot(aes(x = Artist, y = Popularity)) +
  geom_bar(stat = "identity") +
  coord_flip() + labs(title = "popular artists overall", x = "Artists", y = "Popularity")
ggplotly(popular_artists)
```



```
# top artists in based on popularity
```

```
top_10_artists <- spotify2.df %>%
```

```
  group_by(Artist = artists) %>%
```

```
  summarise(No_of_tracks = n(), Popularity = mean(popularity)) %>%
```

```
  filter(No_of_tracks > 1) %>%
```

```
  arrange(desc(Popularity)) %>%
```

```
  top_n(10, wt = Popularity)
```

	Artist <chr>	No_of_tracks <int>	Popularity <dbl>
1	['Bad Bunny', 'Jhay Cortez']	2	91
2	['CJ']	2	86.5
3	['Olivia Rodrigo']	3	85.3
4	['Billie Eilish']	2	85
5	['DJ Khaled', 'Drake']	2	85
6	['Justin Quiles']	2	85
7	['Justin Bieber', 'Khalid']	2	84.5
8	['Miley Cyrus', 'Dua Lipa']	2	84.5
9	['Rauw Alejandro', 'Anuel AA']	2	84.5
10	['ATB', 'Topic', 'A7S']	2	84
11	['Jorge & Mateus']	2	84

We have listed the top artists above with at least 2 songs in the time period so as to exclude the one hit wonders from this analysis.

## DATA MODELING

```
classification_cols =
c("duration_ms","danceability","energy","loudness","speechiness","acousticness"

,"instrumentalness","liveness","valence","tempo","mode","key","explicit","popularity_cat")

# partition

spotify_class.df<-spotify2.df %>% select(classification_cols)

set.seed(1)

train.index <- sample(c(1:dim(spotify_class.df)[1]), dim(spotify_class.df)[1]*0.6)

train.df <- spotify_class.df[train.index, ]

valid.df <- spotify_class.df[-train.index, ]
```

We had split the dataset into training and validation data sets in 60:40 ratio

## DECISION TREE MODELS

```
#decision tree

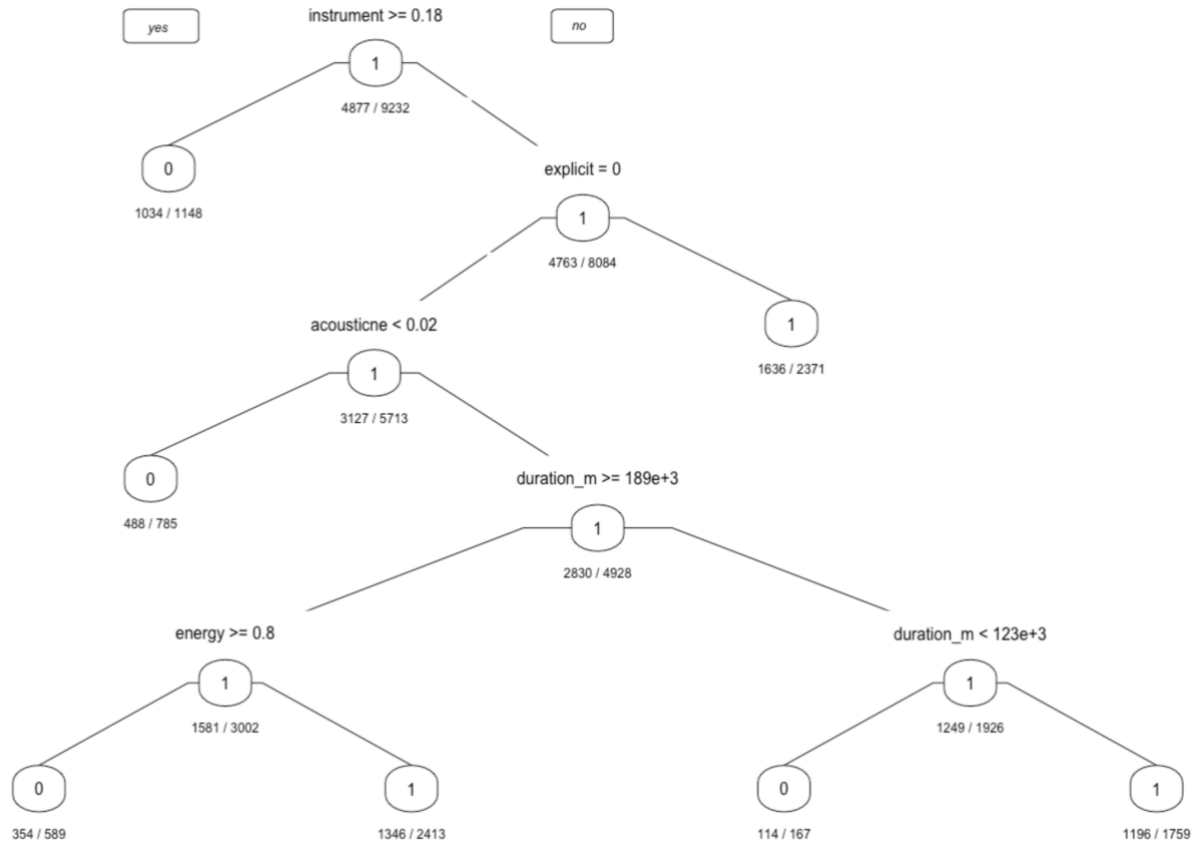
default.ct <- rpart(popularity_cat ~ ., data = train.df ,method = "class")

# plot tree
```

```
prp(default.ct, type = 1, extra = 2, under = TRUE, split.font = 1, varlen = -10)
```

```
# count number of leaves
```

```
length(default.ct$frame$var[default.ct$frame$var == "<leaf>"])
```



There were 7 leaves in the default decision tree.

```
# classify records in the data.
```

```
default.ct.point.pred.train <- predict(default.ct,train.df,type = "class")
```

```
# generate confusion matrix for training data
```

```
confusionMatrix(default.ct.point.pred.train, as.factor(train.df$popularity_cat))
```

```
default.ct.point.pred.valid <- predict(default.ct,valid.df,type = "class")
```

```
confusionMatrix(default.ct.point.pred.valid, as.factor(valid.df$popularity_cat))
```

```
> confusionMatrix(default.ct.point.pred.train, as.factor(train.df$popularity_cat))
Confusion Matrix and Statistics
```

	Reference	
Prediction	0	1
0	1990	699
1	2365	4178

Accuracy : 0.6681  
95% CI : (0.6584, 0.6777)  
No Information Rate : 0.5283  
P-Value [Acc > NIR] : < 0.00000000000000022

```
> confusionMatrix(default.ct.point.pred.valid, as.factor(valid.df$popularity_cat))
Confusion Matrix and Statistics
```

	Reference	
Prediction	0	1
0	1343	472
1	1584	2757

Accuracy : 0.666

#deeper tree

```
deeper.ct <- rpart(popularity_cat ~ ., data = train.df, method = "class", cp = -1, minsplit = 1)
```

```
length(deeper.ct$frame$var[deeper.ct$frame$var == "<leaf>"])
```

There were 1684 leaves in the deeper tree which clearly indicate towards excessive overfitting of data.

### repeat the code for the validation set, and the deeper tree

```
deeper.ct.point.pred.train <- predict(deeper.ct, train.df, type = "class")
```

```
confusionMatrix(deeper.ct.point.pred.train, as.factor(train.df$popularity_cat))
```

```
deeper.ct.point.pred.valid <- predict(deeper.ct, valid.df, type = "class")
```

```
confusionMatrix(deeper.ct.point.pred.valid, as.factor(valid.df$popularity_cat))
```

```
> confusionMatrix(deeper.ct.point.pred.train, as.factor(train.df$popularity_cat))
Confusion Matrix and Statistics
```

	Reference	
Prediction	0	1
0	4350	41
1	5	4836

Accuracy : 0.995

```
> confusionMatrix(deeper.ct.point.pred.valid, as.factor(valid.df$popularity_cat))
Confusion Matrix and Statistics

          Reference
Prediction  0    1
 0  1846 1155
 1  1081 2074

      Accuracy : 0.6368
```

The accuracy was 0.995 for the deeper tree with the training dataset but crashed to 0.6368 with the validation dataset. That clearly suggested the excessive overfitting of data in the original model.

#Pruned tree

```
set.seed(1)
```

```
cv.ct <- rpart(popularity_cat ~ ., data = train.df, method = "class", cp = 0.00001, minsplit = 1, xval = 5)
```

```
printcp(cv.ct)
```

	CP	nsplit	rel error	xerror	xstd
1	0.21125144	0	1.000000	1.00000	0.011014
2	0.02192882	1	0.788749	0.79013	0.010668
3	0.01377727	3	0.744891	0.77015	0.010611
4	0.00711825	6	0.703559	0.72744	0.010475
5	0.00665901	7	0.696441	0.72239	0.010457
6	0.00482204	9	0.683123	0.71504	0.010431
7	0.00401837	10	0.678301	0.71183	0.010419
8	0.00367394	12	0.670264	0.71091	0.010416
9	0.00344432	13	0.666590	0.70402	0.010391
10	0.00275545	14	0.663146	0.70448	0.010393
11	0.00260237	15	0.660390	0.70471	0.010393
12	0.00252583	18	0.652583	0.70448	0.010393
13	0.00241102	19	0.650057	0.70448	0.010393
14	0.00229621	22	0.642250	0.70723	0.010403
15	0.00221967	26	0.633065	0.70677	0.010401
16	0.00206659	30	0.623651	0.70379	0.010390
17	0.00195178	32	0.619518	0.70333	0.010388
18	0.00183697	36	0.610333	0.70333	0.010388
19	0.00172216	37	0.608496	0.70586	0.010398
20	0.00165327	41	0.600459	0.70195	0.010383
21	0.00160735	47	0.590356	0.70333	0.010388
22	0.00153081	50	0.585534	0.70379	0.010390
23	0.00149254	56	0.576349	0.70379	0.010390
24	0.00137773	58	0.573364	0.69874	0.010371
25	0.00126292	63	0.566475	0.69759	0.010367
26	0.00122465	73	0.552928	0.69552	0.010359
27	0.00114811	77	0.547876	0.69552	0.010359
28	0.00107157	89	0.533869	0.69208	0.010346
29	0.00103330	93	0.528817	0.69208	0.010346
30	0.00101033	97	0.524684	0.69208	0.010346
31	0.00099502	106	0.513203	0.69254	0.010347
32	0.00096441	112	0.507233	0.69254	0.010347

From the above diagram, we found that when cp value was 0.00107157, the xerror was minimum but rose after that. The desired cp for the best pruned tree had been found.

```
pruned.ct <- prune(cv.ct, cp = 0.00107157)

length(pruned.ct$frame$var[pruned.ct$frame$var == "<leaf>"])

printcp(pruned.ct)
```

There were 90 leaves in the best pruned tree.

```
# classify records in the training data.
```

```
pruned.ct.point.pred.train <- predict(pruned.ct,train.df,type = "class")
```

```
# generate confusion matrix for training data
```

```
confusionMatrix(pruned.ct.point.pred.train, as.factor(train.df$popularity_cat))
```

```
### repeat the code for the validation set, and the pruned tree
```

```
pruned.ct.point.pred.valid <- predict(pruned.ct,valid.df,type = "class")
```

```
confusionMatrix(pruned.ct.point.pred.valid, as.factor(valid.df$popularity_cat))
```

```
> confusionMatrix(pruned.ct.point.pred.train, as.factor(train.df$popularity_cat))
Confusion Matrix and Statistics
```

	Reference	
Prediction	0	1
0	2796	766
1	1559	4111

Accuracy : 0.7482

```
> confusionMatrix(pruned.ct.point.pred.valid, as.factor(valid.df$popularity_cat))
Confusion Matrix and Statistics
```

	Reference	
Prediction	0	1
0	1670	746
1	1257	2483

Accuracy : 0.6746

```
# random forest
```

```
rf <- randomForest(as.factor(popularity_cat) ~ ., data = train.df, ntree = 500,
```

```
          mtry = 4, nodesize = 5, importance = TRUE)
```

```
## variable importance plot
```

```
varImpPlot(rf, type = 1)
```

```
## confusion matrix
```

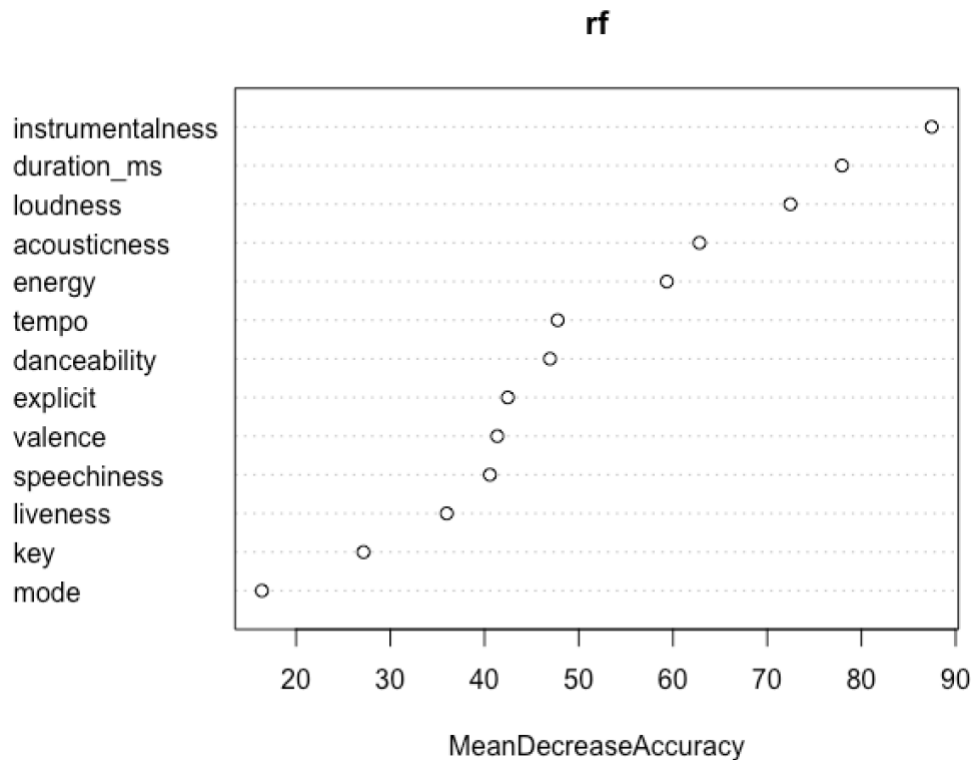
```
rf.pred1 <- predict(rf, train.df)
```



```
confusionMatrix(rf.pred1, as.factor(train.df$popularity_cat))
```

```
rf.pred <- predict(rf, valid.df)
```

```
confusionMatrix(rf.pred, as.factor(valid.df$popularity_cat))
```



```
> confusionMatrix(rf.pred1, as.factor(train.df$popularity_cat))  
Confusion Matrix and Statistics
```

	Reference	
Prediction	0	1
0	4328	38
1	27	4839

Accuracy : 0.993

```
> confusionMatrix(rf.pred, as.factor(valid.df$popularity_cat))  
Confusion Matrix and Statistics
```

	Reference	
Prediction	0	1
0	1691	469
1	1236	2760

Accuracy : 0.723

```

#Boosted tree

library(adabag)

train.df$popularity_cat <- as.factor(train.df$popularity_cat)

set.seed(1)

boost <- boosting(popularity_cat ~ ., data = train.df)

pred1 <- predict(boost, train.df)

confusionMatrix(as.factor(pred1$class), as.factor(train.df$popularity_cat))

pred <- predict(boost, valid.df)

confusionMatrix(as.factor(pred$class), as.factor(valid.df$popularity_cat))

> confusionMatrix(as.factor(pred$class), as.factor(valid.df$popularity_cat))
Confusion Matrix and Statistics

              Reference
Prediction    0      1
0      1560    560
1      1367   2669

              Accuracy : 0.687

> confusionMatrix(as.factor(pred1$class), as.factor(train.df$popularity_cat))
Confusion Matrix and Statistics

              Reference
Prediction    0      1
0      2431    785
1      1924   4092

              Accuracy : 0.7066

```

## LOGISTIC REGRESSION MODELS

```

# regression.

logit.reg <- glm(popularity_cat ~ ., data = train.df, family = "binomial")

options(scipen=999)

summary(logit.reg)

formula(logit.reg)

logit.reg.pred <- predict(logit.reg, valid.df, type = "response")

```

```
logit.reg.pred.classes <- ifelse(logit.reg.pred > 0.5, 1, 0)
```

```
confusionMatrix(as.factor(logit.reg.pred.classes), as.factor(valid.df$popularity_cat))
```

```
> formula(logit.reg)
```

```
popularity_cat ~ duration_ms + danceability + energy + loudness +  
  speechiness + acousticness + instrumentalness + liveness +  
  valence + tempo + mode + key + explicit
```

```
> confusionMatrix(as.factor(logit.reg.pred.classes), as.factor(valid.df$popularity_cat))  
Confusion Matrix and Statistics
```

	Reference	
Prediction	0	1
0	1376	508
1	1551	2721

Accuracy : 0.6655

```
# model selection
```

```
full.logit.reg <- glm(popularity_cat ~ ., data = train.df, family = "binomial")
```

```
empty.logit.reg <- glm(popularity_cat ~ 1, data = train.df, family = "binomial")
```

```
summary(empty.logit.reg)
```

```
backwards = step(full.logit.reg)
```

```
summary(backwards)
```

```
backwards.reg.pred <- predict(backwards, valid.df, type = "response")
```

```
backwards.reg.pred.classes <- ifelse(backwards.reg.pred > 0.5, 1, 0)
```

```
confusionMatrix(as.factor(backwards.reg.pred.classes), as.factor(valid.df$popularity_cat))
```

```
> confusionMatrix(as.factor(backwards.reg.pred.classes), as.factor(valid.df$popularity_cat))  
Confusion Matrix and Statistics
```

	Reference	
Prediction	0	1
0	1367	506
1	1560	2723

Accuracy : 0.6644

```
> formula(backwards)
```

```
popularity_cat ~ duration_ms + danceability + energy + loudness +  
  acousticness + instrumentalness + liveness + explicit
```

```

forwards =
step(empty.logit.reg,scope=list(lower=formula(empty.logit.reg),upper=formula(full.logit.reg)),
direction="forward",trace=0)

formula(forwards)

forwards.reg.pred <- predict(forwards, valid.df, type = "response")

forwards.reg.pred.classes <- ifelse(forwards.reg.pred > 0.5, 1, 0)

confusionMatrix(as.factor(forwards.reg.pred.classes), as.factor(valid.df$popularity_cat))

> confusionMatrix(as.factor(forwards.reg.pred.classes), as.factor(valid.df$popularity_cat))
Confusion Matrix and Statistics

              Reference
Prediction    0      1
0      1367    506
1      1560   2723

              Accuracy : 0.6644

> formula(forwards)
popularity_cat ~ instrumentalness + explicit + loudness + energy +
  acousticness + duration_ms + liveness + danceability

stepwise =
step(empty.logit.reg,scope=list(lower=formula(empty.logit.reg),upper=formula(full.logit.reg)),
direction="both",trace=1)

formula(stepwise)

stepwise.reg.pred <- predict(stepwise, valid.df, type = "response")

stepwise.reg.pred.classes <- ifelse(stepwise.reg.pred > 0.5, 1, 0)

confusionMatrix(as.factor(stepwise.reg.pred.classes), as.factor(valid.df$popularity_cat))

> confusionMatrix(as.factor(stepwise.reg.pred.classes), as.factor(valid.df$popularity_cat))
Confusion Matrix and Statistics

              Reference
Prediction    0      1
0      1367    506
1      1560   2723

              Accuracy : 0.6644

> formula(stepwise)
popularity_cat ~ instrumentalness + explicit + loudness + energy +
  acousticness + duration_ms + liveness + danceability

#Checking for multi collinearity using correlation plots

forcorr_cols=c("duration_ms","danceability","energy","loudness","speechiness","acousticness")

```

```

, "instrumentalness", "liveness", "valence", "tempo", "popularity")

Spotify_Corr.df = spotify2.df %>% select(forcorr_cols)

Cor_spotify = data.frame(cor(Spotify_Corr.df, use = "complete.obs"))


Cor_spotify

plot(Spotify_Corr.df)

plot(Cor_spotify)

ggcorr(Spotify_Corr.df, hjust = 1)

```



From the correlation table, we found that the most highly positively correlated features were loudness and energy with correlation of 0.69 (It was intuitive as higher energy tracks would tend to be louder). Most highly negatively correlated features were acousticness and energy with a value of 0.6.

## FINAL MODEL SELECTION

The accuracies obtained in all the models are given below:

Class	Accuracy obtained on training dataset	Accuracy obtained on validation dataset (from Confusion matrix)
Default decision tree	0.6681	0.666
Deeper tree	0.995	0.6368
Best pruned decision tree	0.7482	0.6746
Random forest	0.993	0.7238
Boosted tree	0.7066	0.687
Default logistic regression	-	0.6655
Backward selection regression	-	0.6644
Forward selection regression	-	0.6644
Stepwise selection regression	-	0.6644

- **Random forest model offers the highest accuracy on the validation dataset and is selected as the model for prediction of track popularity**
- Interesting takeaway: The selection algorithms didn't improve the performance of the default regression model
- Key features which influence popularity of a track based on Decision tree, Randomforest and regression were:
  - instrumentalness, explicit, energy, acousticness, loudness
  - Duration\_ms (This was surprising since we do not usually consider the length of a track while listening to music)

## CLUSTERING WITH K MEANS

```
cluster_cols=c("id","duration_ms","danceability","energy","loudness","speechiness","acousticness",  
              ,"instrumentalness","liveness","valence","tempo")
```

```
spotify_cluster.df <- spotify2.df %>% select(cluster_cols)
```

```
row.names(spotify_cluster.df) <- spotify_cluster.df[,1]
```

```
spotify_cluster.df <- spotify_cluster.df[,-1]
```

```
View(spotify_cluster.df)
```

```
summary(spotify_cluster.df)
```

```
spotify_cluster.df.norm <- sapply(spotify_cluster.df, scale)
```

```
row.names(spotify_cluster.df.norm) <- row.names(spotify_cluster.df)
```

```
# Plotting elbow chart


k.max <- 50

data <- spotify_cluster.df.norm

wss <- sapply(1:k.max,
              function(k){kmeans(data, k, nstart=50,iter.max = 15 )$tot.withinss})

wss

plot(1:k.max, wss,
     type="b", pch = 19, frame = FALSE,
     xlab="Number of clusters K",
     ylab="Total within-clusters sum of squares")
```



From the above elbow plot, we concluded that 8 clusters would be appropriate for performing the k means algorithm.

```
km <- kmeans(spotify_cluster.df.norm,8)

#Plots for g

# plot an empty scatter plot
```

```


plot(c(0), xaxt = 'n', ylab = "", type = "l",
     ylim = c(min(km$centers), max(km$centers)), xlim = c(0, 10))

# label x-axes
axis(1, at = c(1:10), labels = names(spotify_cluster.df))

# plot centroids
for (i in c(1:8))
  lines(km$centers[i,], lty = i, lwd = 3, col = ifelse(i %in% c(1, 3, 5),
                                                         "black", "dark grey"))

# name clusters
text(x = 0.3, y = km$centers[, 1], labels = paste("Cluster", c(1:8)))

```



Cluster Name	Interpretation of characteristics with respect to other clusters
Cluster 1	High duration_ms and low instrumentalness (which means spoken words are more prominent compared to instruments. This cluster might contain audiobooks and podcasts.
Cluster 2	Low acousticness but very high instrumentalness. This cluster might contain rock songs.
Cluster 3	These stores sell the highest Original jeans with second highest sales in Fashion jeans and average sales in other jean types. The reasons for the extraordinarily high



	sales should be studied to find out if the learnings can help boost the sales in other clusters.
Cluster 4	Highest speechiness with high danceability. This cluster might have vocal focused tracks with less instruments but good danceability.
Cluster 5	Low duration_ms, lowest loudness and highest instrumentality. This might contain theme music like Indiana Jones music
Cluster 6	Highest duration, high speechiness. This might contain Marathon tracks of artists and Yearbook mix
Cluster 7	Highest liveness. This might contain Live recordings like concert
Cluster 8	Lowest duration but high energy and high instrumentality. This might contain rock music

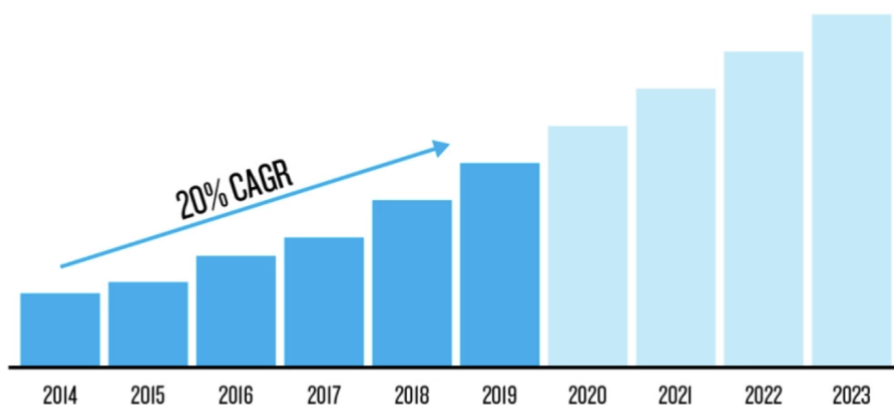
- Valence and tempo have good variation across clusters

## KEY TAKEAWAYS

- For popularity prediction, random forest had the highest accuracy on the validation dataset and hence the model can be used for future predictions of track popularity
- From the logistic regression model, we inferred that acoustics tracks with high loudness and danceability with explicit lyrics can become popular songs on Spotify
- Duration\_ms also play a major role in the popularity. This would not have been the case a few years but a rise of popularity of audio books, talk shows and podcasts might have increased the significance of duration of the tracks in popularity as all these new listening areas have higher duration compared to regular music tracks. Nielson research shared below shows CAGR growth of about 20% for podcast listeners in the US.

## PODCAST AUDIENCE GROWTH RATE

The U.S. podcast audience could double by 2023



Source: Nielsen Podcast Listener Buying Power Database

Copyright © 2020 The Nielsen Company (US), LLC. All Rights Reserved.

Source: <https://www.nielsen.com/insights/2020/podcast-content-is-growing-audio-engagement/>